

CSE 6220 Introduction to High Performance Computing
Spring 2014
Programming Assignment 2
Due March 31

Write a parallel program to sort integers in parallel by doing a straightforward parallelization of the sequential quicksort algorithm. The algorithm is described here at a high-level and you should fill in the remaining implementation details.

The input to the algorithm is a file containing the number of integers to be sorted followed by the integers themselves. The output of the program is the sorted list of the integers. You may structure the program such that one processor reads the input file and sends relevant data to other processors. Note that we are doing this for convenience sake, and because we do not have a machine that supports truly parallel I/O.

Let n denote the number of integers and p denote the number of processors. The integers are equally distributed to the processors: Each processor has an array A holding either $\lceil \frac{n}{p} \rceil$ or $\lfloor \frac{n}{p} \rfloor$ integers. The algorithm is recursive: At some stage during the algorithm, suppose there is a communicator of size q on which m integers should be sorted and that they are distributed in the usual manner. If $q = 1$, then any serial sorting algorithm can be used to sort the input. You may either write your code, copy from a book or web site if available, or even use the Unix sort program.

If $q > 1$, all processors in the communicator use the same random number generator to generate a random number between 0 and $m - 1$, say k . The processor that has the k^{th} integer (called pivot from here on) broadcasts it to all processors in the communicator. Each processor goes through its integer array and partitions it into two subarrays – one containing integers less than or equal to the pivot and another containing integers greater than the pivot. Using an *Allgather* operation, the number of integers in each of the two subarrays on all the processors are gathered on every processor. Compute the total number of integers less than or equal to the pivot (say m') and the total number of integers greater than the pivot (say m''). Partition the p processors to the two subproblems of sorting m' integers and sorting m'' integers respectively, by allocating processors in proportion to the problem sizes (make sure you do not allocate zero processors to a problem). Using the gathered information on the sizes of the subarrays on all processors, each processor can compute where to send its data and from where to receive its data. Use an *Alltoall* communication to perform the data transfer. Create two new communicators corresponding to the two partitions. Recursively sort within each partition.

The program will be graded on correctness and usage of the required parallel programming features. You should also use good programming style and comment your program so that it is readable and understandable.

Experimentally evaluate the performance of your program. For instance, fix a problem size, vary the number of processors, and plot the run-time as a function of the number of processors. Fix the

number of processors and see what happens as the problem size is varied. Try to come up with some important observations on your program. What is the minimum problem size per processor to get good parallel efficiency? I am not asking that specific questions be answered but you are expected to think on your own, experiment and show the conclusions along with the evidence that led you to reach the conclusions.

The programming assignment is to be done by groups of two students each. It is important that you strictly adhere to the input format described in the programming assignment. No matter how you decide to share the work amongst yourselves, each student is expected to have full knowledge of the submitted program. To submit the programming assignment upload a `tar.gz` file containing:

1. A text file containing the name of people on the team and their contributions.
2. Source code provide a README to compile and run your program.
3. Description of design decisions and analyses in pdf format.

On jinx you should use the command `tar czf assignment2.tar.gz assignment2` in order to compress the directory containing your code and report prior to upload to t-square. Projects that are provided in formats other than a `tar.gz` file containing plain text and pdf files will not be graded. There is a sample project attached. This has all of the necessary components for grading named appropriately. Please modify the contents of each file for your work and submit on t-square.